




deeplearning.ai

Batch Normalization

Why does
Batch Norm work?

Learning on shifting input distribution



Cat

Non-Cat

$y = 1$

$y = 0$




$y = 1$

$y = 0$



Why this is a problem with neural networks?



Batch Norm as regularization

- Each mini-batch is scaled by the mean/variance computed on just that mini-batch.
- This adds some noise to the values $z^{[l]}$ within that minibatch. So similar to dropout, it adds some noise to each hidden layer's activations.
- This has a slight regularization effect.